

## Session 8 Flow (90–120 min)

1.

### Revisit Agents (10 min)

- Quick recap of **Session 6 & 7**:
  - Agents = reasoning + action.
  - Tools = how agents act.
  - MCP = standard “USB port” for tool integration.
- Slide reminder: **ReAct loop** (Plan → Act → Observe → Reflect).

2.

### Intro to Agentic RAG (15 min)

- RAG alone = knowledge from docs.
- Agents alone = action with tools.
- **Agentic RAG = knowledge + action in one loop.**
- Show the diagram (the one you liked):
  - Query → Agent → Plan → (RAG retriever OR MCP tool) → Reflect → Answer.
- Emphasize **enterprise fit**: GitLab, Confluence, Jira, internal DBs.

3.

### MCP Refresher (10 min)

- Why MCP is critical: security, portability, reusability.
- Tools are just APIs, MCP makes them **plug-and-play**.
- Example: GitLab MCP server with `search_code`, `list_mrs`, `get_file`.

4.

### Architecture Walkthrough (15 min)

- Present the **Mermaid diagrams** (System, Sequence, Security).
- Highlight flow:
  - **RAG path** = Chroma + Sentence-Transformers.
  - **MCP path** = Files MCP + GitLab MCP.
  - **LLM-only path** for fallback.
- Show **security guardrails**: allowlists, read-only tokens, logging.

## 5.

### Live Demo (30 min)

Run the upgraded **Agentic RAG Enterprise Project**:

- **RAG example:**  
“Summarize the onboarding policy.”  
→ See retrieval chunks in UI.
- **GitLab example:**  
“Search for login in project 123 (main).”  
→ Trace shows MCP tool call → GitLab API.
- **Files MCP example:**  
“List files under ./rag-service/data.”
- **LLM-only example:**  
“What is the capital of France?”  
→ Goes direct to LLM fallback.
- Emphasize **timeline trace in UI** (plan, tool call, results).

## 6.

### Wrap-Up (10 min)

- Agentic RAG = **bridge** between knowledge and action.
  - Demos showed:
    - Chroma retriever (knowledge).
    - GitLab MCP (enterprise system).
    - Unified orchestrator + UI trace.
  - Next step in roadmap: **Orchestration frameworks** (LangGraph, Prefect, etc.).
-



## Tips to Position Yourself as Expert

- Use enterprise language: “governance,” “allowlists,” “audit logs,” “scalability.”
- Drop references to real-world use: GitLab MR review bots, Merchant onboarding, Knowledge copilots.
- During demo, pause after each query to explain **why the orchestrator routed it that way**.
- End with: “You’ve now seen how we can plug real enterprise systems into AI copilots in a safe and standardized way.”

### Governance

- **Scenario:** Explaining *why MCP is safer* than ad-hoc tool integrations.
- **Where to say it:** Right after introducing MCP servers.
- **Phrase:**

*“MCP gives us governance — central control over which tools agents can call, with predictable schemas and monitoring hooks.”*

---



### Allowlists

- **Scenario:** Presenting the **GitLab MCP server**.
- **Where to say it:** When showing .env with ALLOWED\_PROJECTS=123,456.
- **Phrase:**

*“We don’t just let agents loose on the whole GitLab — we enforce project-level allowlists so only approved repos are accessible.”*

---



### Audit Logs

- **Scenario:** Describing **enterprise observability & compliance**.
- **Where to say it:** While explaining your **Security & Ops Mermaid diagram**.
- **Phrase:**

*“Every MCP tool call — what tool, what args, which user — is logged. These audit logs can feed into SIEM systems, ensuring compliance and traceability.”*

---




## Scalability

- **Scenario:** Wrapping up the demo with forward-looking vision.
- **Where to say it:** In the **Wrap-Up / Q&A**.
- **Phrase:**

*“Because the architecture is modular — RAG retrievers, MCP servers, orchestrator — it scales horizontally. We can add Jira MCP, Confluence MCP, or new retrievers without breaking the system.”*

---

 By sprinkling these terms in those **specific points of your session**, you’ll sound not only like someone who understands the tech, but also like an **enterprise architect** who’s thought through governance, compliance, and scale.

# Session 8 – Presenter Cheat Sheet

## 1. Revisit Agents

- “Agents are not just answering, they **plan, act, and reflect**.”
- “Tools are how agents act — and with MCP, tools become plug-and-play.”

## 2. Intro to Agentic RAG

- “RAG = knowledge, Agents = action. **Agentic RAG = both, in one loop.**”
- “The orchestrator decides: do I retrieve knowledge, call a tool, or just answer?”

## 3. MCP Refresher

- “MCP is the governance layer — like a USB standard for AI tools.”
- “It enforces **schemas, allowlists, and logs** — so we can trust what tools are being called.”

## 4. Architecture

- “Here’s the flow: User → Orchestrator → either RAG, MCP, or LLM-only.”
- “Notice how GitLab MCP is read-only, allowlisted, and logged — that’s enterprise compliance built in.”

## 5. Live Demo (say before each query)

- “Now let’s try a RAG query — retrieving policy docs.”
- “Here’s a GitLab query — search for ‘login’ in project 123.”
- “This one is a local file MCP query.”
- “Finally, a general knowledge query goes to LLM-only.”

(While the UI shows the **timeline**!)

- “See how the orchestrator shows its plan, the tool call, the raw results, and the final synthesis.”

## 6. Wrap-Up

- “We’ve now seen how **knowledge** + **action** are combined in a safe, governed way.”
- “Governance via MCP, allowlists on projects, audit logs on tool calls.”
- “This architecture is scalable — we can plug in Jira, Confluence, Vault, or any enterprise system the same way.”

Here are **two clear reasons** (keep them sharp for your session):

---

1

### Uncontrolled External Execution

- **Why it’s a concern:**

MCP servers expose tools that agents can call at runtime.

If not tightly governed, an LLM could trigger **file access, code execution, or API calls** beyond intended scope.

- **How security sees it:**

*“This opens potential vectors for **data leakage** or **privilege escalation**, unless every tool is strictly sandboxed and allowlisted.”*

---

2

### Compliance & Audit Gaps

- **Why it’s a concern:**

MCP tooling is still new. Enterprises don't yet have **standardized audit trails, RBAC policies, or SIEM connectors** for MCP traffic.

- **How security sees it:**

*“Without mature logging, monitoring, and approval workflows, MCP tool calls can't meet compliance frameworks like SOC2, ISO27001, PCI-DSS.”*

---

👉 You can then add:

*“That's why in today's demo, MCP is scoped to **read-only, allowlisted GitLab projects and local file paths**, with output caps and logging — the very guardrails security would expect before approving wider MCP adoption.”*

---

## 🔴 Why MCP is restricted today (your org's security stance)

1. **Uncontrolled external execution**
    - MCP lets LLMs trigger tools at runtime.
    - Without strict sandboxing, this risks **data leakage** or **privilege escalation**.
  2. **Compliance & audit gaps**
    - MCP is new; lacks mature **audit trails, RBAC, SIEM connectors**.
    - Hard to prove compliance for **SOC2, ISO27001, PCI-DSS**.
- 

## 🟢 How MCP can become enterprise-safe

1. **Governance & allowlists**
    - Limit tools to **read-only, allowlisted systems** (e.g., GitLab projects).
    - Enforce **schema validation, rate limits, and output caps**.
    - Wrap every MCP call with **audit logging** → SIEM.
  2. **Controlled deployment**
    - Run MCP servers **inside the secure VPC/Kubernetes cluster**, behind internal auth.
    - Add **RBAC policies** so only approved agent workflows can call sensitive tools.
    - Integrate with **existing enterprise monitoring** (Splunk, Azure Monitor, etc.).
- 

✅ That way, you position yourself as both:

- **Realistic** (why it's blocked today).
- **Forward-looking** (how it can be adopted safely).

### 💡 Analogy for MCP Security

*“Think of MCP like a USB port for AI. Just as you can plug in any USB device to your laptop — a keyboard, a flash drive, or even something malicious — MCP lets an AI agent plug into any tool. IT doesn't block USBs because they're bad, they block them until there are policies: device control, encryption, monitoring. MCP is in the same stage — powerful, but it needs those guardrails before enterprises allow it.”*

---

That makes it **memorable** and shows you understand both **the risk** and **the governance path forward**.