

## Session 4 – Adaptive & Agentic RAG (Goals)

- Show a minimal ReAct-style planner that picks tools (vector, SQL, synth).
- Return final answer + execution trace (for observability).
- Discuss memory, tool routing, retry, and guardrails.

# Agentic Flow

- 1) LLM plans steps (tools list).
- 2) Retrieve vector context; optionally run SQL tool based on question.
- 3) Synthesize grounded answer with sources; show trace.

# Enterprise Considerations

- Audit logs, prompt versioning, PII redaction, safety filters.
- Latency/cost vs accuracy knobs; caching and reranking.
- Path to production: tracing (OpenTelemetry), feature flags, evals.