

## Session 1 – Transcript (RAG Foundations)

- Welcome: Starting RAG journey today.
- Definition: RAG = Retrieval-Augmented Generation – grounding LLM answers with external knowledge.
- Why needed: LLM cutoff, hallucinations, missing enterprise context.
- Demo: Chunk docs → Embed → Store in vector DB → Query retrieves relevant chunks.
- Example: 'What is valid address proof?' → returns context from RBI doc.
- Wrap-up: Naive RAG gives quick wins for FAQs and support.